
CSE 573: Artificial Intelligence

Hanna Hajishirzi
Markov Decision Processes

slides adapted from
Dan Klein, Pieter Abbeel ai.berkeley.edu
And Dan Weld, Luke Zettelmoyer




Review and Outline

- Adversarial Games
 - Minimax search
 - α - β search
 - Evaluation functions
 - Multi-player, non-0-sum
- Stochastic Games
 - Expectimax
 - Markov Decision Processes
 - Reinforcement Learning




Agents vs. Environment

- An **agent** is an entity that *perceives* and *acts*.
- A **rational agent** selects actions that *maximize its utility function*.




Deterministic **vs. stochastic**
Fully observable **vs.** partially observable

Non-Deterministic Search




Example: Grid World

- A maze-like problem
 - The agent lives in a grid
 - Walls block the agent's path
- Noisy movement: actions do not always go as planned
 - 80% of the time, the action North takes the agent North (if there is no wall there)
 - 10% of the time, North takes the agent West; 10% East
 - If there is a wall in the direction the agent would have been taken, the agent stays put
- The agent receives rewards each time step
 - Small "living" reward each step (can be negative)
 - Big rewards come at the end (good or bad)
- Goal: maximize sum of rewards




Grid World Actions

Deterministic Grid World



Stochastic Grid World



Markov Decision Processes

- An MDP is defined by:


- A set of states $s \in S$
- A set of actions $a \in A$
- A transition function $T(s, a, s')$
 - Probability that a from s leads to s' , i.e., $P(s' | s, a)$
 - Also called the model or the dynamics

$T(s_{11}, E, \dots)$

\dots
 $T(s_{31}, N, s_{11}) = 0$

\dots
 $T(s_{31}, N, s_{32}) = 0.8$
 $T(s_{31}, N, s_{21}) = 0.1$
 $T(s_{31}, N, s_{41}) = 0.1$
 \dots

T is a Big Table!
 $11 \times 4 \times 11 = 484$ entries




For now, we give this as input to the agent

Markov Decision Processes

- An MDP is defined by:

- A set of states $s \in S$
- A set of actions $a \in A$
- A transition function $T(s, a, s')$
 - Probability that a from s leads to s' , i.e., $P(s' | s, a)$
 - Also called the model or the dynamics
- A reward function $R(s, a, s')$
 - Sometimes just $R(s)$ or $R(s')$

$$P(a|s)$$



$R(s_{32}, N, s_{33}) = -0.01$

\dots

$R(s_{32}, N, s_{42}) = -1.01$

\dots

$R(s_{33}, E, s_{43}) = 0.99$


Cost of breathing

R is also a Big Table!

For now, we also give this to the agent

Markov Decision Processes


- An MDP is defined by:
 - A **set of states** $s \in S$
 - A **set of actions** $a \in A$
 - A **transition function** $T(s, a, s')$
 - Probability that a from s leads to s' , i.e., $P(s' | s, a)$
 - Also called the model or the dynamics
 - A **reward function** $R(s, a, s')$
 - Sometimes just $R(s)$ or $R(s')$
 - A **start state**
 - Maybe a **terminal state**
- MDPs are non-deterministic search problems
 - One way to solve them is with expectimax search
 - We'll have a new tool soon



What is Markov about MDPs?

- “Markov” generally means that given the present state, the future and the past are independent
- For Markov decision processes, “Markov” means action outcomes depend only on the current state

$$P(S_{t+1} = s' | S_t = s_t, A_t = a_t, S_{t-1} = s_{t-1}, A_{t-1}, \dots, S_0 = s_0)$$

= 

$$P(S_{t+1} = s' | S_t = s_t, A_t = a_t)$$




Andrey Markov
(1856-1922)

- This is just like search, where the successor function could only depend on the current state (not the history)





Policies

- In deterministic single-agent search problems, we wanted an optimal **plan**, or sequence of actions, from start to a goal
- For MDPs, we want an optimal policy $\pi^*: S \rightarrow A$
 - A policy π gives an action for each state
 - An optimal policy is one that maximizes expected utility if followed
 - An explicit policy defines a reflex agent




Optimal policy when $R(s, a, s') = -0.4$ for all non-terminals s

Optimal Policies




Example: Racing




Example: Racing

- A robot car wants to travel far, quickly
- Three states: Cool, Warm, Overheated
- Two actions: *Slow*, *Fast*
- Going faster gets double reward




Racing Search Tree




MDP Search Trees

- Each MDP state projects an expectimax-like search tree




Utilities of Sequences



Utilities of Sequences

- What preferences should an agent have over reward sequences?
- More or less? $\underline{[1, 2, 2]}$ or $[2, \underline{3}, \underline{4}]$
- Now or later? $\underline{[0, 0, 1]}$ or $[1, \underline{0}, \underline{0}]$



Discounting

- It's reasonable to maximize the sum of rewards
- It's also reasonable to prefer rewards now to rewards later
- One solution: values of rewards decay exponentially



1

Worth Now



γ



γ^2

Worth In Two Steps

Discounting

- How to discount?


- Each time we descend a level,
we multiply in the discount once

- Why discount?

- Think of it as a gamma chance of
ending the process at every step
- Also helps our algorithms
converge


- Example: discount of 0.5

- $U([1,2,3]) = 1 * 1 + 0.5 * 2 + 0.25 * 3$
- $U([1,2,3]) < U([3,2,\cancel{1}])$ $\gamma^2 \approx ?$




Quiz: Discounting

- Given:



- Actions: East, West, and Exit (only available in exit states a, e)
- Transitions: deterministic

- Quiz 1: For $\gamma = 1$, what is the optimal policy?



- Quiz 2: For $\gamma = 0.1$, what is the optimal policy?



- Quiz 3: For which γ are West and East equally good when in state d?

$$1\gamma=10\gamma^3$$

$$\gamma = 10\gamma^3$$

Infinite Utilities?

- Problem: What if the game lasts forever? Do we get infinite rewards?

- Solutions:


- Finite horizon: (similar to depth-limited search)

- Terminate episodes after a fixed T steps (e.g. life)
 - Policy π depends on time left

- Discounting: use $0 < \gamma < 1$

$$U([r_0, \dots, r_\infty]) = \sum_{t=0}^{\infty} \gamma^t r_t \leq R_{\max} / (1 - \gamma)$$


- Smaller γ means smaller “horizon” – shorter term focus
 - Absorbing state: guarantee that for every policy, a terminal state will eventually be reached (like “overheated” for racing)



Recap: Defining MDPs

- Markov decision processes:

- Set of states \underline{S}
- Start state $\underline{s_0}$
- Set of actions \underline{A}
- Transitions $P(\underline{s' | s,a})$ (or $T(s,a,s')$)
- Rewards $R(s,a,s')$ (and discount γ)



- MDP quantities so far:

- Policy = Choice of action for each state
- Utility = sum of (discounted) rewards

Solving MDPs

