


CSE 573: Artificial Intelligence



Hanna Hajishirzi
HMMs Inference, Particle Filters

slides adapted from
Dan Klein, Pieter Abbeel ai.berkeley.edu
And Dan Weld, Luke Zettelmoyer




Recap: Reasoning Over Time

- Markov models




- Hidden Markov models



X	E	P
rain	umbrella	0.9
rain	no umbrella	0.1
sun	umbrella	0.2
sun	no umbrella	0.8

Conditional Independence

- HMMs have two important independence properties:
 - Markov hidden process: future depends on past via the present
 - Current observation independent of all else given current state



- Does this mean that evidence variables are guaranteed to be independent?
 - [No, they tend to correlate by the hidden state]

Real HMM Examples


- Robot tracking:
 - Observations are range readings (continuous)
 - States are positions on a map (continuous)
- Speech recognition HMMs:
 - Observations are acoustic signals (continuous valued)
 - States are specific positions in specific words (so, tens of thousands)
- Machine translation HMMs:
 - Observations are words (tens of thousands)
 - States are translation options

Filtering / Monitoring

- Filtering, or monitoring, is the task of tracking the distribution $B_t(X) = P_t(X_t | e_1, \dots, e_t)$ (the belief state) over time
- We start with $B_1(X)$ in an initial setting, usually uniform
- As time passes, or we get observations, we update $B(X)$
- The Kalman filter was invented in the 60's and first implemented as a method of trajectory estimation for the Apollo program

Example: Robot Localization


Example from
Michael Pfeiffer




Sensor model: can read in which directions there is a wall,
never more than 1 mistake

Motion model: may not execute action with small prob.


Example: Robot Localization




Lighter grey: was possible to get the reading, but less likely b/c required 1 mistake




Example: Robot Localization




Example: Robot Localization



Example: Robot Localization



Example: Robot Localization



Inference: Find State Given Evidence


- We are given evidence at each time and want to know

$$B_t(X) = P(X_t | e_{1:t})$$


e₁ e₂ . . . e_t

- Idea: start with $P(X_1)$ and derive B_t in terms of B_{t-1}
 - equivalently, derive B_{t+1} in terms of B_t

Inference: Base Cases




$$P(\underline{\underline{X_1}} | \underline{\underline{e_1}}) = \frac{P(e_1 | X_1) P(X_1)}{Z}$$



$$\begin{aligned} & P(\underline{\underline{X_2}}) \\ &= \sum_{x_1} P(x_1, \underline{\underline{X_2}}) \\ &= \sum_{x_1} P(X_2 | x_1) P(x_1) \end{aligned}$$

Inference: Base Cases



$$\begin{aligned} P(x_2) &= \sum_{x_1} P(x_1, x_2) \\ &= \sum_{x_1} P(x_1)P(x_2|x_1) \end{aligned}$$

Passage of Time

- Assume we have current belief $P(X \mid \text{evidence to date})$

$$B(X_t) = P(X_t | e_{1:t})$$


- Then, after one time step passes:

$$\begin{aligned} P(X_{t+1} | e_{1:t}) &= \sum_{x_t} P(X_{t+1}, x_t | e_{1:t}) \\ &= \sum_{x_t} P(X_{t+1} | x_t, e_{1:t}) P(x_t | e_{1:t}) \\ &= \sum_{x_t} P(X_{t+1} | x_t) P(x_t | e_{1:t}) \end{aligned}$$



Basic idea: beliefs get “pushed” through the transitions

- With the “B” notation, we have to be careful about what time step t the belief is about, and what evidence it includes



- Or compactly:

$$B'(X_{t+1}) = \sum_{x_t} P(X' | x_t) B(x_t)$$

Example: Passage of Time

- As time passes, uncertainty “accumulates”

(Transition model: ghosts usually go clockwise)

<0.01	<0.01	<0.01	<0.01	<0.01	<0.01
<0.01	<0.01	<0.01	<0.01	<0.01	<0.01
<0.01	<0.01	1.00	<0.01	<0.01	<0.01
<0.01	<0.01	<0.01	<0.01	<0.01	<0.01



T = 1

<0.01	<0.01	<0.01	<0.01	<0.01	<0.01
<0.01	<0.01	0.06	<0.01	<0.01	<0.01
<0.01	0.76	0.06	0.06	<0.01	<0.01
<0.01	<0.01	0.06	<0.01	<0.01	<0.01


T = 2

0.05	0.01	0.05	<0.01	<0.01	<0.01
0.02	0.14	0.11	0.35	<0.01	<0.01
0.07	0.03	0.05	<0.01	0.03	<0.01
0.03	0.03	<0.01	<0.01	<0.01	<0.01

T = 5



Inference: Base Cases



$$P(X_1|e_1)$$

ꝝ

$$P(x_1|e_1) = P(x_1, e_1)/P(e_1)$$

$$\propto_{X_1} P(x_1, e_1)$$

$$= P(x_1)P(e_1|x_1)$$

Observation

- Assume we have current belief $P(X \mid \text{previous evidence})$:

$$B'(X_{t+1}) = P(X_{t+1} | e_{1:t})$$

- Then, after evidence comes in:


$$\begin{aligned} P(X_{t+1} | e_{1:t+1}) &= P(X_{t+1}, e_{t+1} | e_{1:t}) / P(e_{t+1} | e_{1:t}) \\ &\propto_{X_{t+1}} P(X_{t+1}, e_{t+1} | e_{1:t}) \end{aligned}$$

$$= P(e_{t+1} | e_{1:t}, X_{t+1}) P(X_{t+1} | e_{1:t})$$

$$= P(e_{t+1} | X_{t+1}) \underbrace{P(X_{t+1} | e_{1:t})}_{\leftarrow \rightarrow}$$

- Or, compactly:

$$\underbrace{B(X_{t+1})}_{\leftarrow \rightarrow} \propto_{X_{t+1}} P(e_{t+1} | X_{t+1}) \underbrace{B'(X_{t+1})}_{\leftarrow \rightarrow}$$



- Basic idea: beliefs “reweighted” by likelihood of evidence
- Unlike passage of time, we have to renormalize

Example: Observation

- As we get observations, beliefs get reweighted, uncertainty “decreases”



0.05	0.01	0.05	<0.01	<0.01	<0.01
0.02	0.14	0.11	0.35	<0.01	<0.01
0.07	0.03	0.05	<0.01	0.03	<0.01
0.03	0.03	<0.01	<0.01	<0.01	<0.01


Before observation




<0.01	<0.01	<0.01	<0.01	0.02	<0.01
<0.01	<0.01	<0.01	0.83	0.02	<0.01
<0.01	<0.01	0.11	<0.01	<0.01	<0.01
<0.01	<0.01	<0.01	<0.01	<0.01	<0.01

After observation

$$B(X) \propto P(e|X)B'(X)$$




Pacman – Sonar (P4)



Recap: HMMs

- HMMs have two important independence properties:
 - Markov hidden process: future depends on past via the present
 - Current observation independent of all else given current state



- Does this mean that evidence variables are guaranteed to be independent?
 - [No, they tend to correlate by the hidden state]




Filtering: $P(X_t \mid \text{evidence}_{1:t})$

Elapse time: compute $P(X_t \mid e_{1:t-1})$

$$P(x_t | e_{1:t-1}) = \sum_{x_{t-1}} P(x_{t-1} | e_{1:t-1}) \cdot P(x_t | x_{t-1})$$

Observe: compute $P(X_t | e_{1:t})$

$$P(x_t | e_{1:t}) \propto P(x_t | e_{1:t-1}) \cdot P(e_t | x_t)$$



Belief: $\langle P(\text{rain}), P(\text{sun}) \rangle$


$P(X_1)$	$<0.5, 0.5>$	<i>Prior on X_1</i>
----------	--------------	----------------------------------

$P(X_1 \mid E_1 = \text{umbrella})$	$<0.82, 0.18>$	<i>Observe</i>
-------------------------------------	----------------	----------------

$P(X_2 \mid E_1 = \text{umbrella})$	$<0.63, 0.37>$	<i>Elapse time</i>
-------------------------------------	----------------	--------------------

$P(X_2 \mid E_1 = \text{umbrella}, E_2 = \text{umbrella})$	$<0.88, 0.12>$	<i>Observe</i>
--	----------------	----------------


Example: Weather HMM



Approximate Inference


- Sometimes $|X|$ is too big for exact inference
 - $|X|$ may be too big to even store $B(X)$
 - E.g. when X is continuous
 - $|X|^2$ may be too big to do updates
- Solution: approximate inference by sampling
- How robot localization works in practice

Approximate Inference: Sampling



Sampling

- Sampling is a lot like repeated simulation
 - Predicting the weather, basketball games, ...
- Basic idea
 - Draw N samples from a sampling distribution S
 - Compute an approximate probability
- Why sample?
 - Learning: get samples from a distribution you don't know
 - Inference: getting a sample is faster than computing the right answer



Sampling

- Sampling from given distribution

- Step 1: Get sample u from uniform distribution over $[0, 1]$


- E.g. `random()` in python

0.61

- Step 2: Convert this sample u into an outcome for the given distribution by having each target outcome associated with a sub-interval of $[0,1)$ with sub-interval size equal to probability of the outcome

- Example

C	P(C)
red	0.6
green	0.1
blue	0.3




$0 \leq u < 0.6, \rightarrow C = \text{red}$


$0.6 \leq u < 0.7, \rightarrow C = \text{green}$

$0.7 \leq u < 1, \rightarrow C = \text{blue}$

- If `random()` returns $u = 0.83$, then our sample is $C = \text{blue}$
 - E.g, after sampling 8 times:




Particle Filtering




Particle Filtering

- Filtering: approximate solution
- Sometimes $|X|$ is too big to use exact inference
 - $|X|$ may be too big to even store $B(X)$
 - E.g. X is continuous
- Solution: approximate inference
 - Track samples of X , not all values
 - Samples are called particles
 - Time per step is linear in the number of samples
 - But: number needed may be large
 - In memory: list of particles, not states
- This is how robot localization works in practice
- Particle is just new name for sample



Representation: Particles

- Our representation of $P(X)$ is now a list of N particles (samples)
 - Generally, $N \ll |X|$
 - Storing map from X to counts would defeat the point
- $P(x)$ approximated by number of particles with value x
 - So, many x may have $P(x) = 0!$
 - More particles, more accuracy
- For now, all particles have a weight of 1



Particle Filtering: Elapse Time

- Each particle is moved by sampling its next position from the transition model

$$\underline{x}' = \text{sample}(P(X'|x))$$


- Samples' frequencies reflect the transition probabilities
- Here, most samples move clockwise, but some move in another direction or stay in place
- This captures the passage of time
 - If enough samples, close to exact values before and after (consistent)

Particles:

- (3,3)
- (2,3)
- (3,3)
- (3,2)
- (3,3)
- (3,2)
- (1,2)
- (3,3)
- (3,3)
- (2,3)

Particles:

- (3,2)
- (2,3)
- (3,2)
- (3,1)
- (3,3)
- (3,2)
- (1,3)
- (2,3)
- (3,2)
- (2,2)



Particle Filtering: Observe

- Slightly trickier:

- Don't sample observation, fix it
- Downweight samples based on the evidence


$$w(x) = P(e|x)$$

$$B(X) \propto \underbrace{P(e|X)}_{\text{Evidence}} B'(X)$$

- As before, the probabilities don't sum to one, since all have been downweighted (in fact they now sum to (N times) an approximation of $P(e)$)


Particles:

(3,2)
(2,3)
(3,2)
(3,1)
(3,3)
(3,2)
(1,3)
(2,3)
(3,2)
(2,2)



Particles:

(3,2) w=.9
(2,3) w=.2
(3,2) w=.9
(3,1) w=.4
(3,3) w=.4
(3,2) w=.9
(1,3) w=.1
(2,3) w=.2
(3,2) w=.9
(2,2) w=.4




Particle Filtering: Resample

- Rather than tracking weighted samples, we resample
- N times, we choose from our weighted sample distribution (i.e. draw with replacement)
- This is equivalent to renormalizing the distribution
- Now the update is complete for this time step, continue with the next one


Particles:

(3,2) w=.9
(2,3) w=.2
(3,2) w=.9
(3,1) w=.4
(3,3) w=.4
(3,2) w=.9
(1,3) w=.1
(2,3) w=.2
(3,2) w=.9
(2,2) w=.4




(New) Particles:

(3,2)
(2,2)
(3,2)
(2,3)
(3,3)
(3,2)
(1,3)
(2,3)
(3,2)
(3,2)



Recap: HMMs

- HMMs have two important independence properties:
 - Markov hidden process: future depends on past via the present
 - Current observation independent of all else given current state



- Does this mean that evidence variables are guaranteed to be independent?
 - [No, they tend to correlate by the hidden state]




Filtering: $P(X_t \mid \text{evidence}_{1:t})$

Elapse time: compute $P(X_t \mid e_{1:t-1})$

$$P(x_t | e_{1:t-1}) = \sum_{x_{t-1}} P(x_{t-1} | e_{1:t-1}) \cdot P(x_t | x_{t-1})$$

Observe: compute $P(X_t | e_{1:t})$

$$P(x_t | e_{1:t}) \propto P(x_t | e_{1:t-1}) \cdot P(e_t | x_t)$$



Belief: $\langle P(\text{rain}), P(\text{sun}) \rangle$


$P(X_1)$	$<0.5, 0.5>$	<i>Prior on X_1</i>
----------	--------------	----------------------------------

$P(X_1 \mid E_1 = \text{umbrella})$	$<0.82, 0.18>$	<i>Observe</i>
-------------------------------------	----------------	----------------


$P(X_2 \mid E_1 = \text{umbrella})$	$<0.63, 0.37>$	<i>Elapse time</i>
-------------------------------------	----------------	--------------------

$P(X_2 \mid E_1 = \text{umbrella}, E_2 = \text{umbrella})$	$<0.88, 0.12>$	<i>Observe</i>
--	----------------	----------------

Video (Markov Model)




Video (HMM)




Particle Filtering

- Filtering: approximate solution
- Sometimes $|X|$ is too big to use exact inference
 - $|X|$ may be too big to even store $B(X)$
 - E.g. X is continuous
- Solution: approximate inference
 - Track samples of X , not all values
 - Samples are called particles
 - Time per step is linear in the number of samples
 - But: number needed may be large
 - In memory: list of particles, not states
- This is how robot localization works in practice
- Particle is just new name for sample



Recap: Particle Filtering


- Particles: track samples of states rather than an explicit distribution




$$x' = \text{sample}(P(X'|x))$$

$$w(x) = P(e|x)$$

Video of Demo – Moderate Number of Particles




Video of Demo – Huge Number of Particles




Which Algorithm?

Particle filter, uniform initial beliefs, 25 particles




Which Algorithm?

Exact filter, uniform initial beliefs



Which Algorithm?




Particle filter, uniform initial beliefs, 300 particles



Robot Localization

- In robot localization:

- We know the map, but not the robot's position
- Observations may be vectors of range finder readings
- State space and readings are typically continuous (works basically like a very fine grid) and so we cannot store $B(X)$
- Particle filtering is a main technique




Particle Filter Localization (Sonar)



The image shows a simulation environment for particle filter localization. A grey rectangular arena contains several obstacles represented by red clusters of points. A green dot at the bottom left represents the robot's current position. A blue box in the bottom-left corner displays the number "40000". Overlaid on the image is the text "Global localization with sonar sensors" in large, bold, black font.

**Global localization with
sonar sensors**

Particle Filter Localization (Laser)




Most Likely Explanation



HMMs: MLE Queries

- HMMs defined by

- States X
- Observations E
- Initial distribution: $P(X_1)$
- Transitions: $P(X|X_{-1})$
- Emissions: $P(E|X)$




- New query: most likely explanation:

$$\arg \max_{x_{1:t}} P(x_{1:t}|e_{1:t})$$

- New method: the Viterbi algorithm


State Trellis

- State trellis: graph of states and transitions over time



- Each arc represents some transition $x_{t-1} \rightarrow x_t$
- Each arc has weight $P(x_t|x_{t-1})P(e_t|x_t)$
- Each path is a sequence of states
- The product of weights on a path is that sequence's probability along with the evidence
- Forward algorithm computes sums of paths, Viterbi computes best paths

Forward / Viterbi Algorithms



Forward Algorithm (Sum)

$$f_t[x_t] = P(x_t, e_{1:t})$$

$$= P(e_t|x_t) \sum_{x_{t-1}} P(x_t|x_{t-1}) f_{t-1}[x_{t-1}]$$

Viterbi Algorithm (Max)

$$m_t[x_t] = \max_{x_{1:t-1}} P(x_{1:t-1}, x_t, e_{1:t})$$

$$= P(e_t|x_t) \max_{x_{t-1}} P(x_t|x_{t-1}) m_{t-1}[x_{t-1}]$$